

March 22, 2023

DSML: Computer Vision.

CNNs : Operations under the hood .

Class starts
@ 9:05 pm.



What normal people see
when they walk on street



What Computer Vision
folks see



WHO WOULD WIN?



STATE OF THE ART
NEURAL NETWORK



ONE NOISY BOI

Recap:

I] Operation

- ① Conv2D (valid)
- ② Conv2D (same)
- ③ Conv2D (full)
- ④ Conv2D (strided)
- ⑤ Conv2D (dilated)
- ⑥ MaxPool2D
- ⑦ AvgPool2D

Input sizes

<u>Kernel</u>	<u>Image</u>
---------------	--------------

$a \times b$	$m \times n$
--------------	--------------

$a \times b$	$m \times n$
--------------	--------------

$a \times b$	$m \times n$
--------------	--------------

$a \times b$	$m \times n, k \times k$
--------------	--------------------------

$a \times b$	$m \times n, k \times k$
--------------	--------------------------

$a \times b$	$m \times n$
--------------	--------------

$a \times b$	$m \times n$
--------------	--------------

Output sizes

$m - a + 1 \times n - b + 1$	
------------------------------	--

$m \times n$	
--------------	--

$m + a - 1 \times n + b - 1$	
------------------------------	--

$\frac{m}{k} \times \frac{n}{k}$	
----------------------------------	--

$m \times n$	
--------------	--

$\frac{m}{a} \times \frac{n}{b}$	
----------------------------------	--

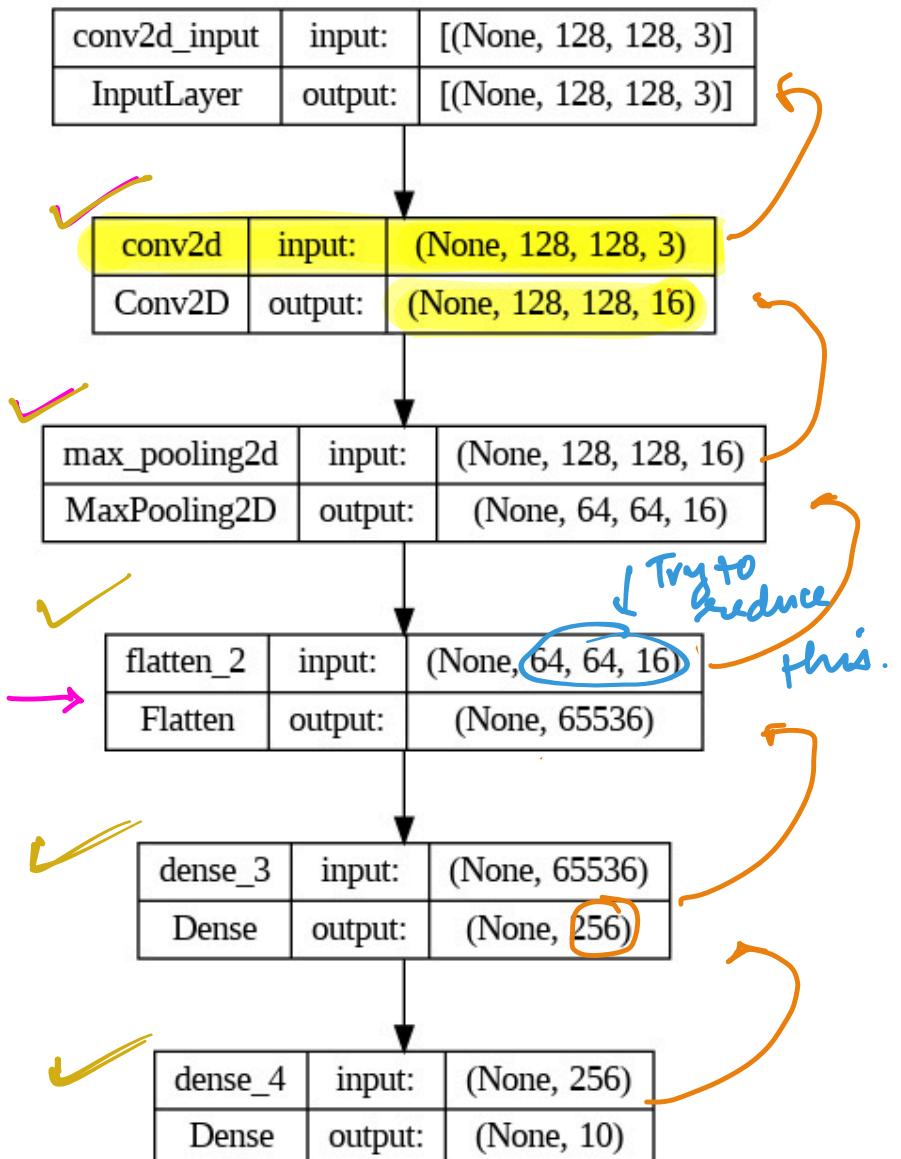
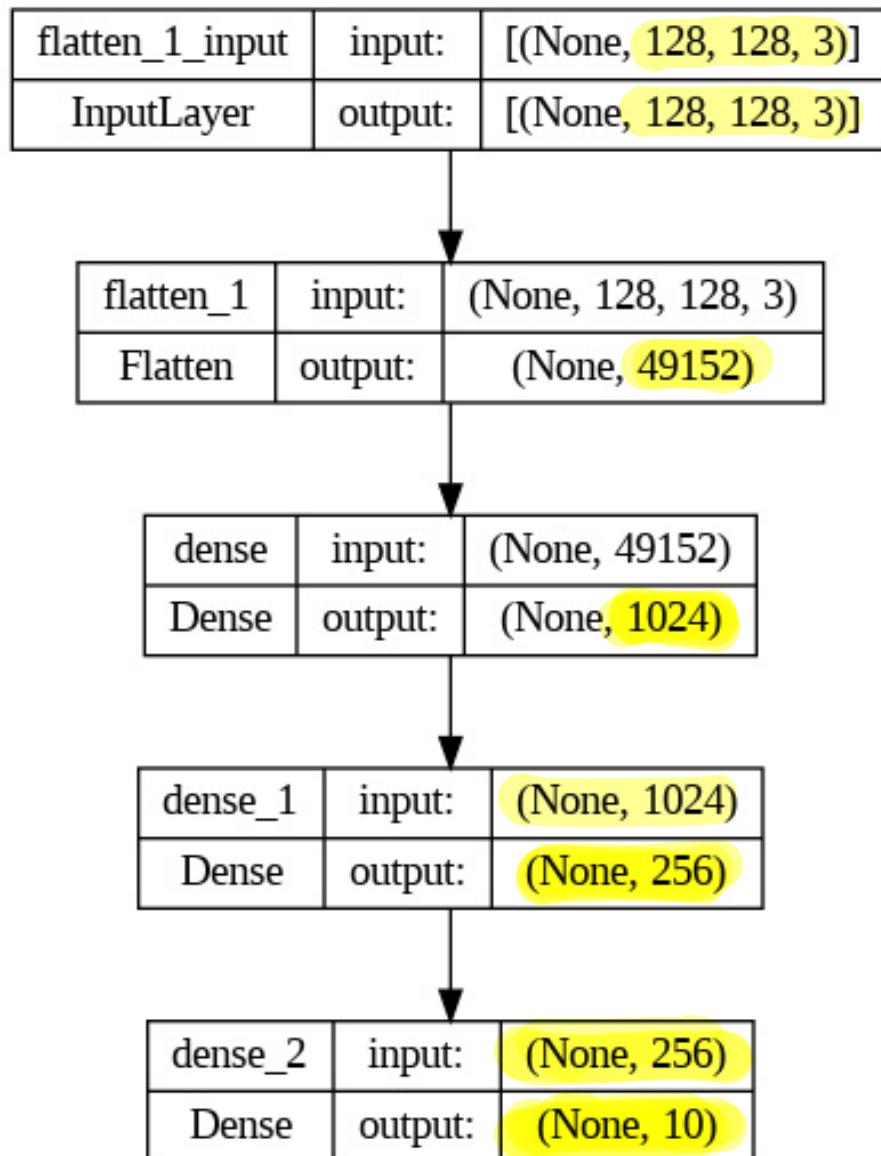
$\frac{m}{a} \times \frac{n}{b}$	
----------------------------------	--

II] Apparel classification Case study:

ANN : $\approx 40\%$.

CNN : $\approx 50\%$.

Architectures :

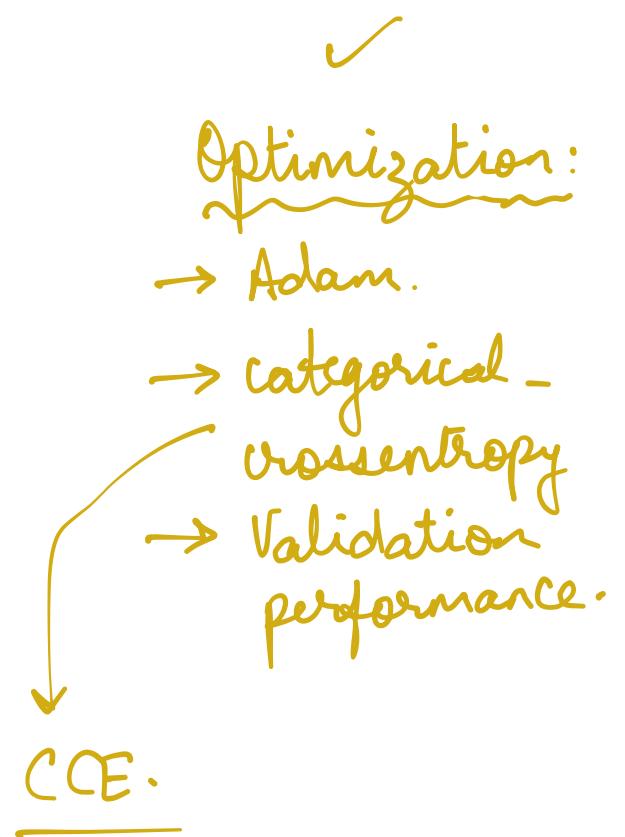


Overview: Major Components

Data: ✓

- $128 \times 128 \times 3$ Images.
- 10 categories.

Model: ✓



① Regularization (Dropout, Batchnorm)

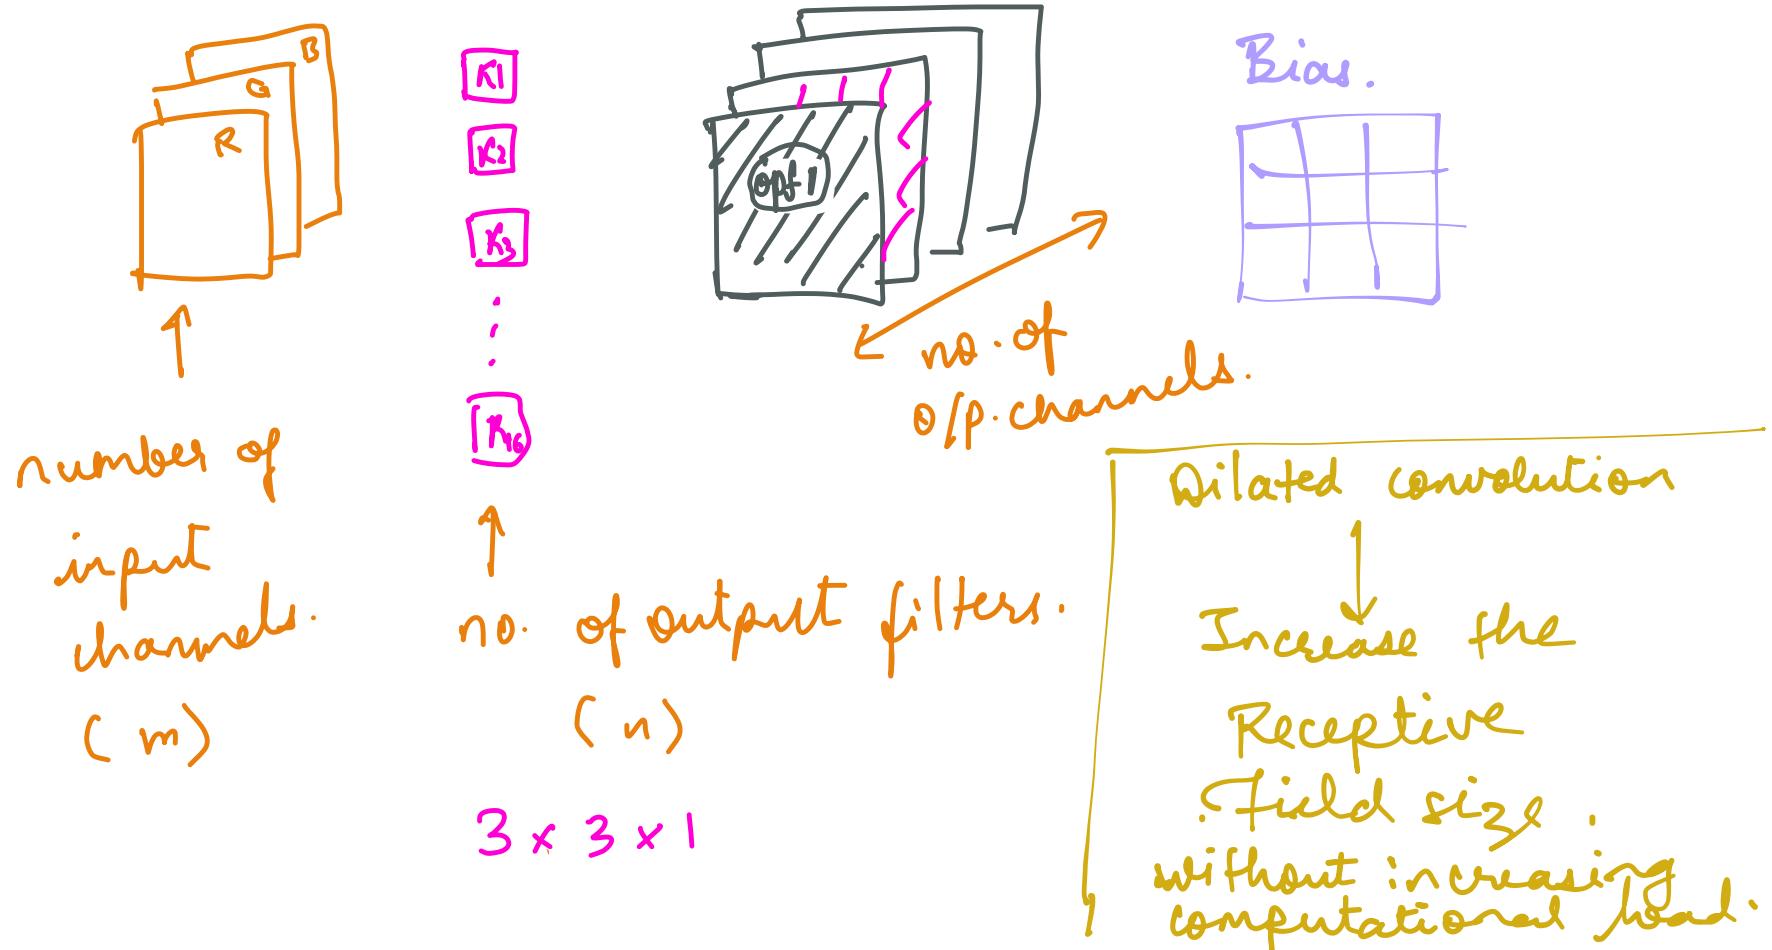
② More data.

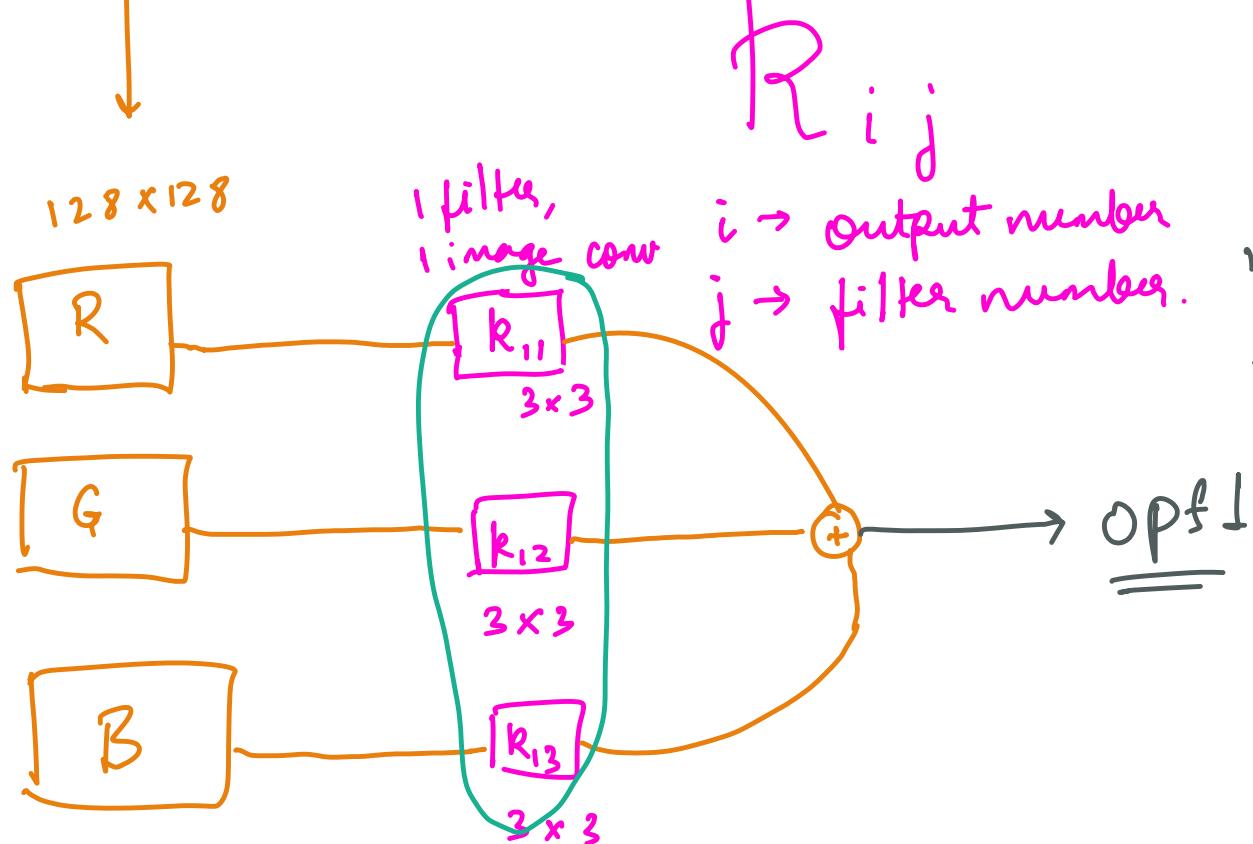
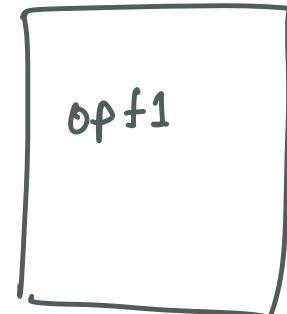
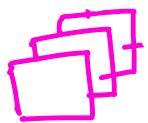
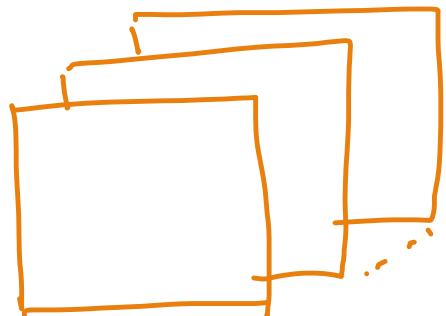
③ Reduce parameters. (Decrease layers etc.)

④ Changing opts, batch, learning rate etc, hyperparams, early stopping.

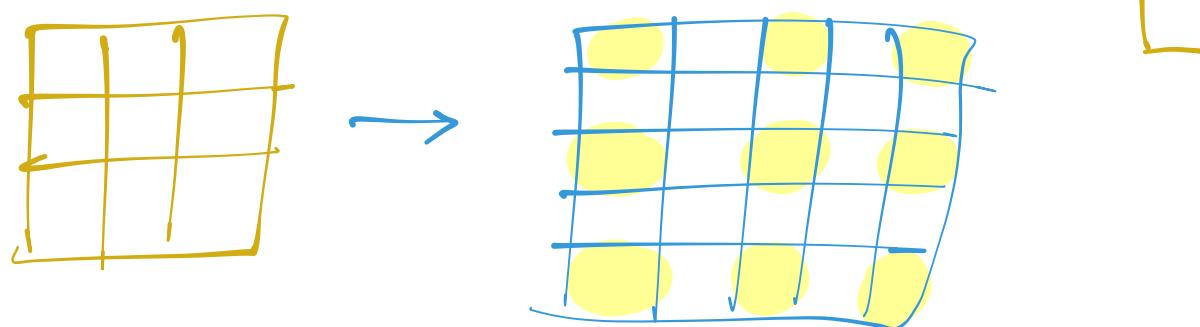
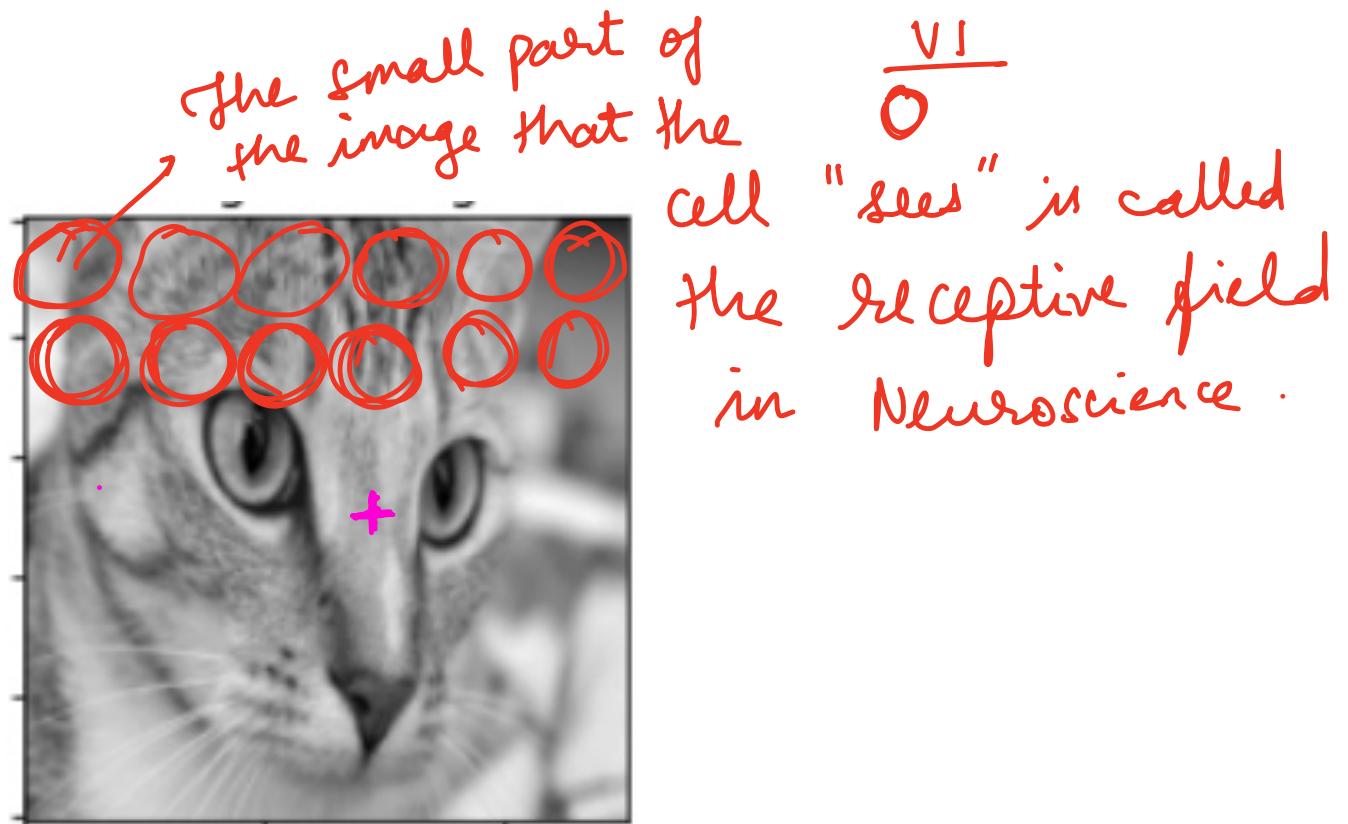
The Conv2D block:

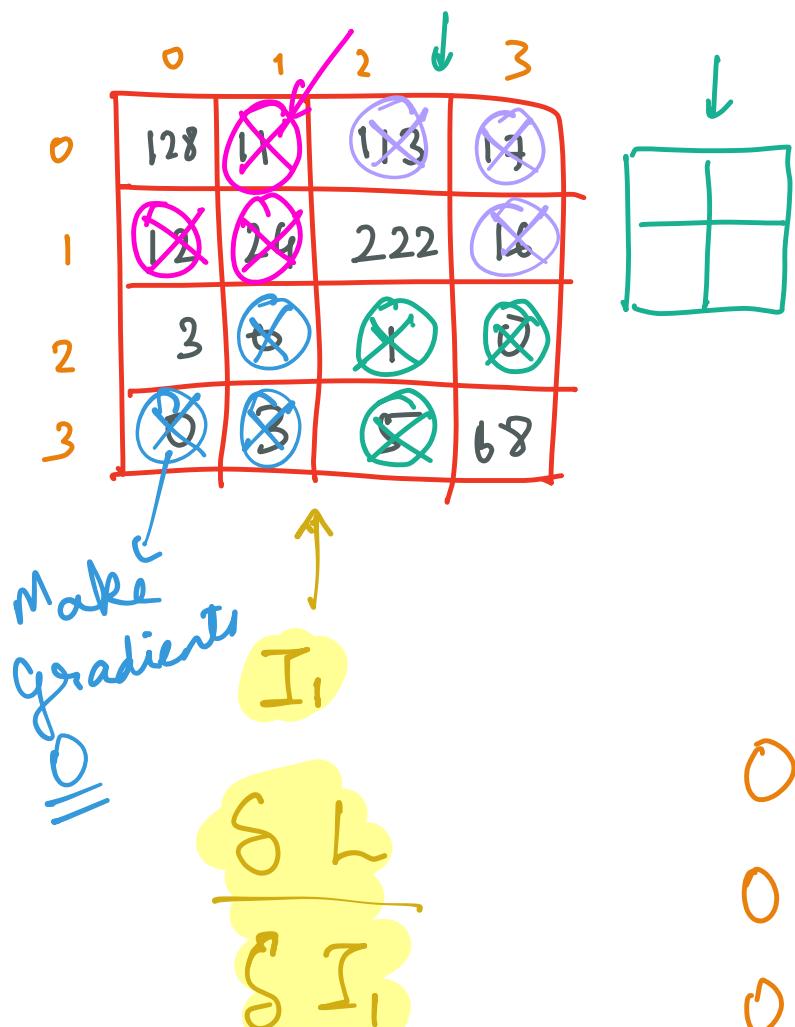
- Please try implementing different conv2D operations from scratch.
- from scipy.signal import convolve2D



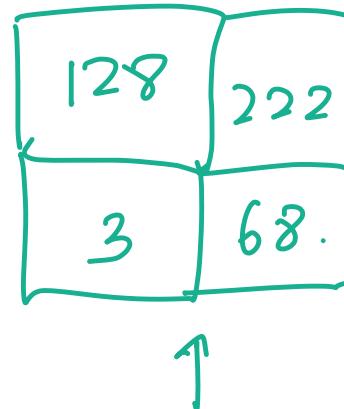


To calculate this, we need 3 filters if we have 3 input channels.



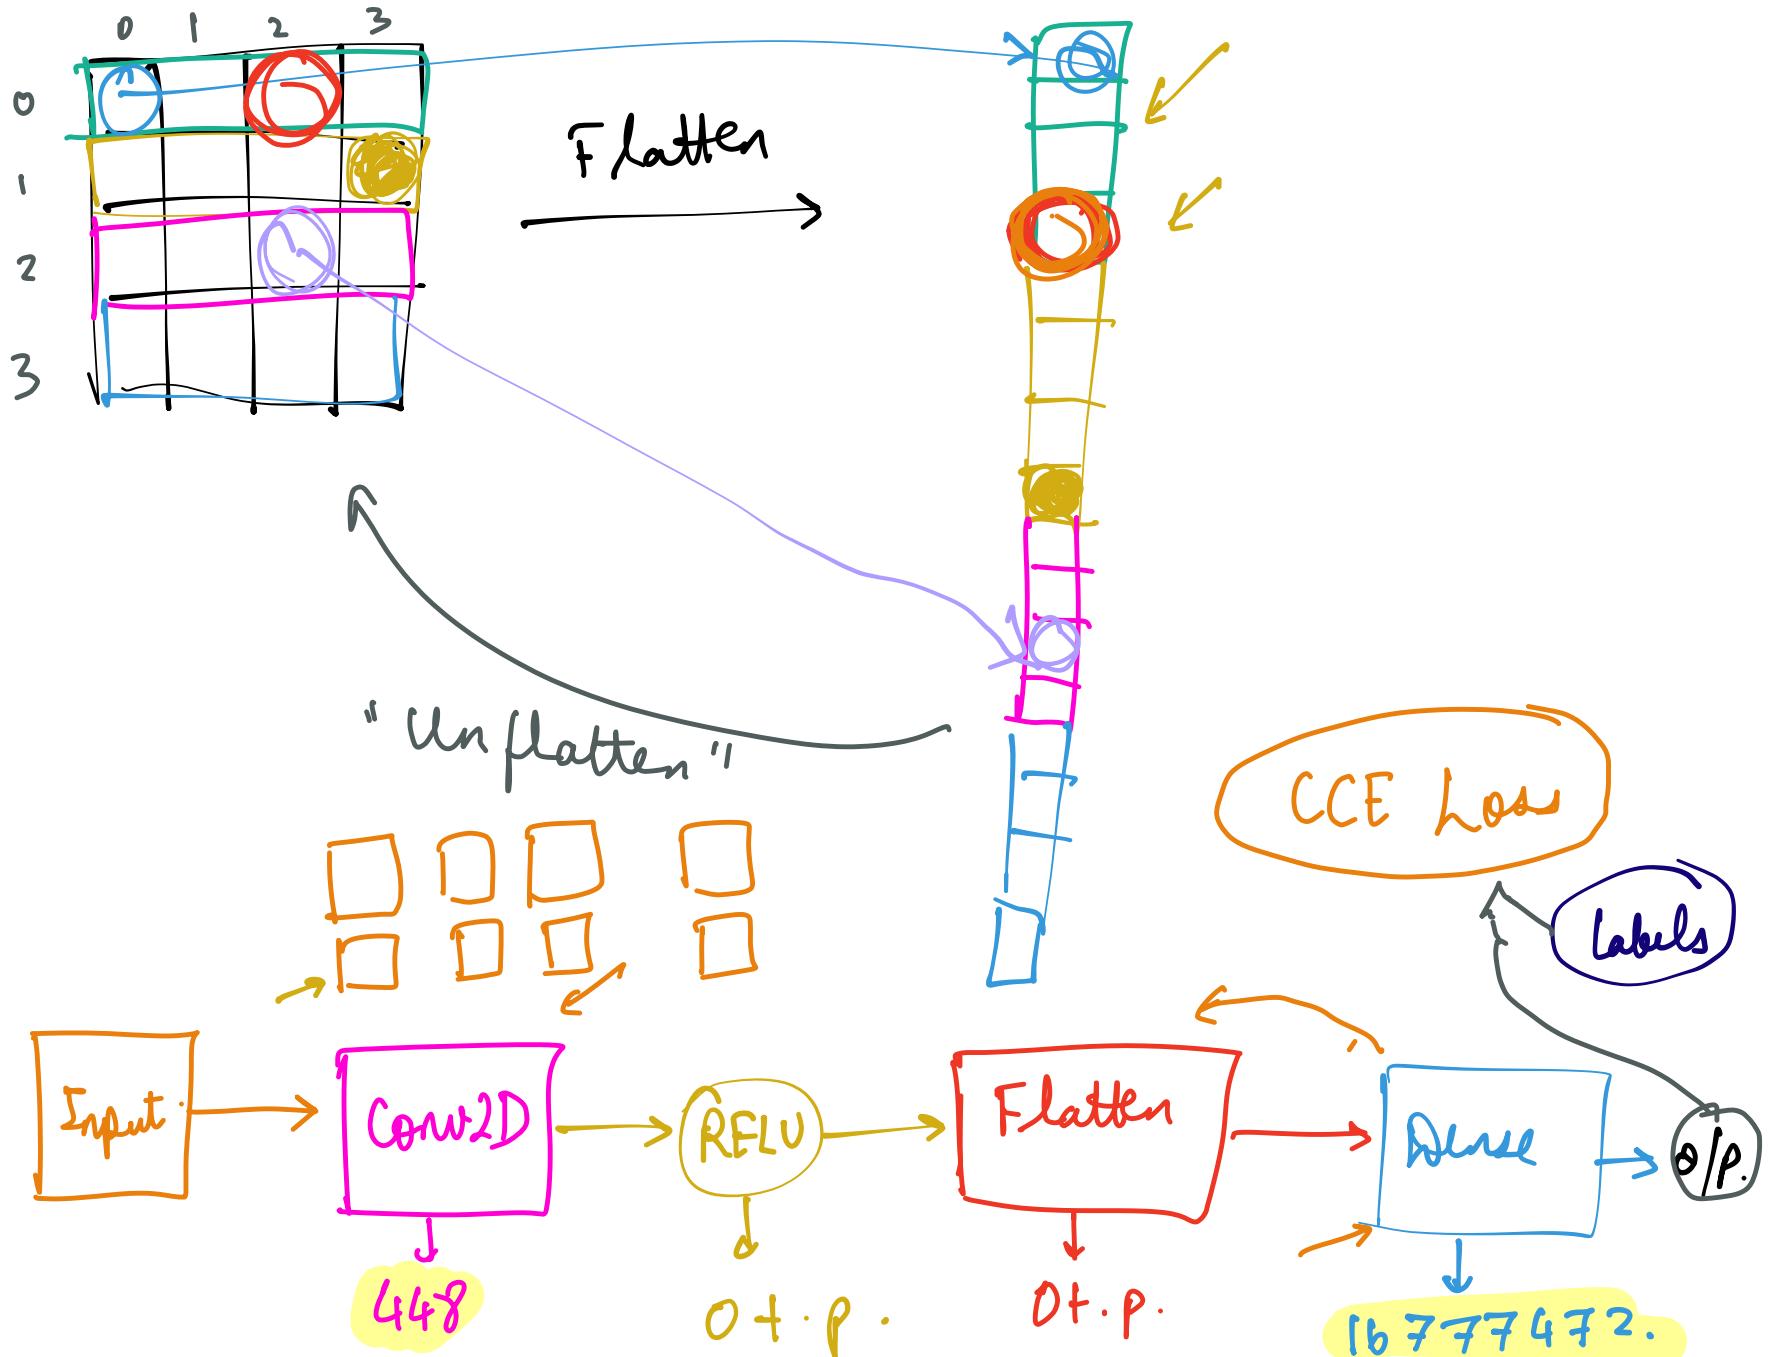


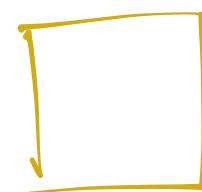
Backprop for max pool



$$\begin{aligned}
 O_1[0,0] &= I_1[0,0] \\
 O_1[0,1] &= I_1[1,2] \\
 O_1[1,0] &= I_1[2,0] \\
 O_1[1,1] &= I_1[3,3].
 \end{aligned}$$

For all indices of I_1 , which are not picked,
We assume the gradient contribution to be 0.





$n \times n \times 3$

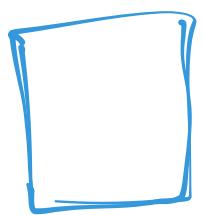


$3 \times 3 \times k$

$n \times n \times k$

Objective :

Reduce
this.



$\square_{1 \times 1}$

dimensions
of kernels.

- ✓ ① Maxpool.
- ✓ ② Strided conv.
with stride 2.

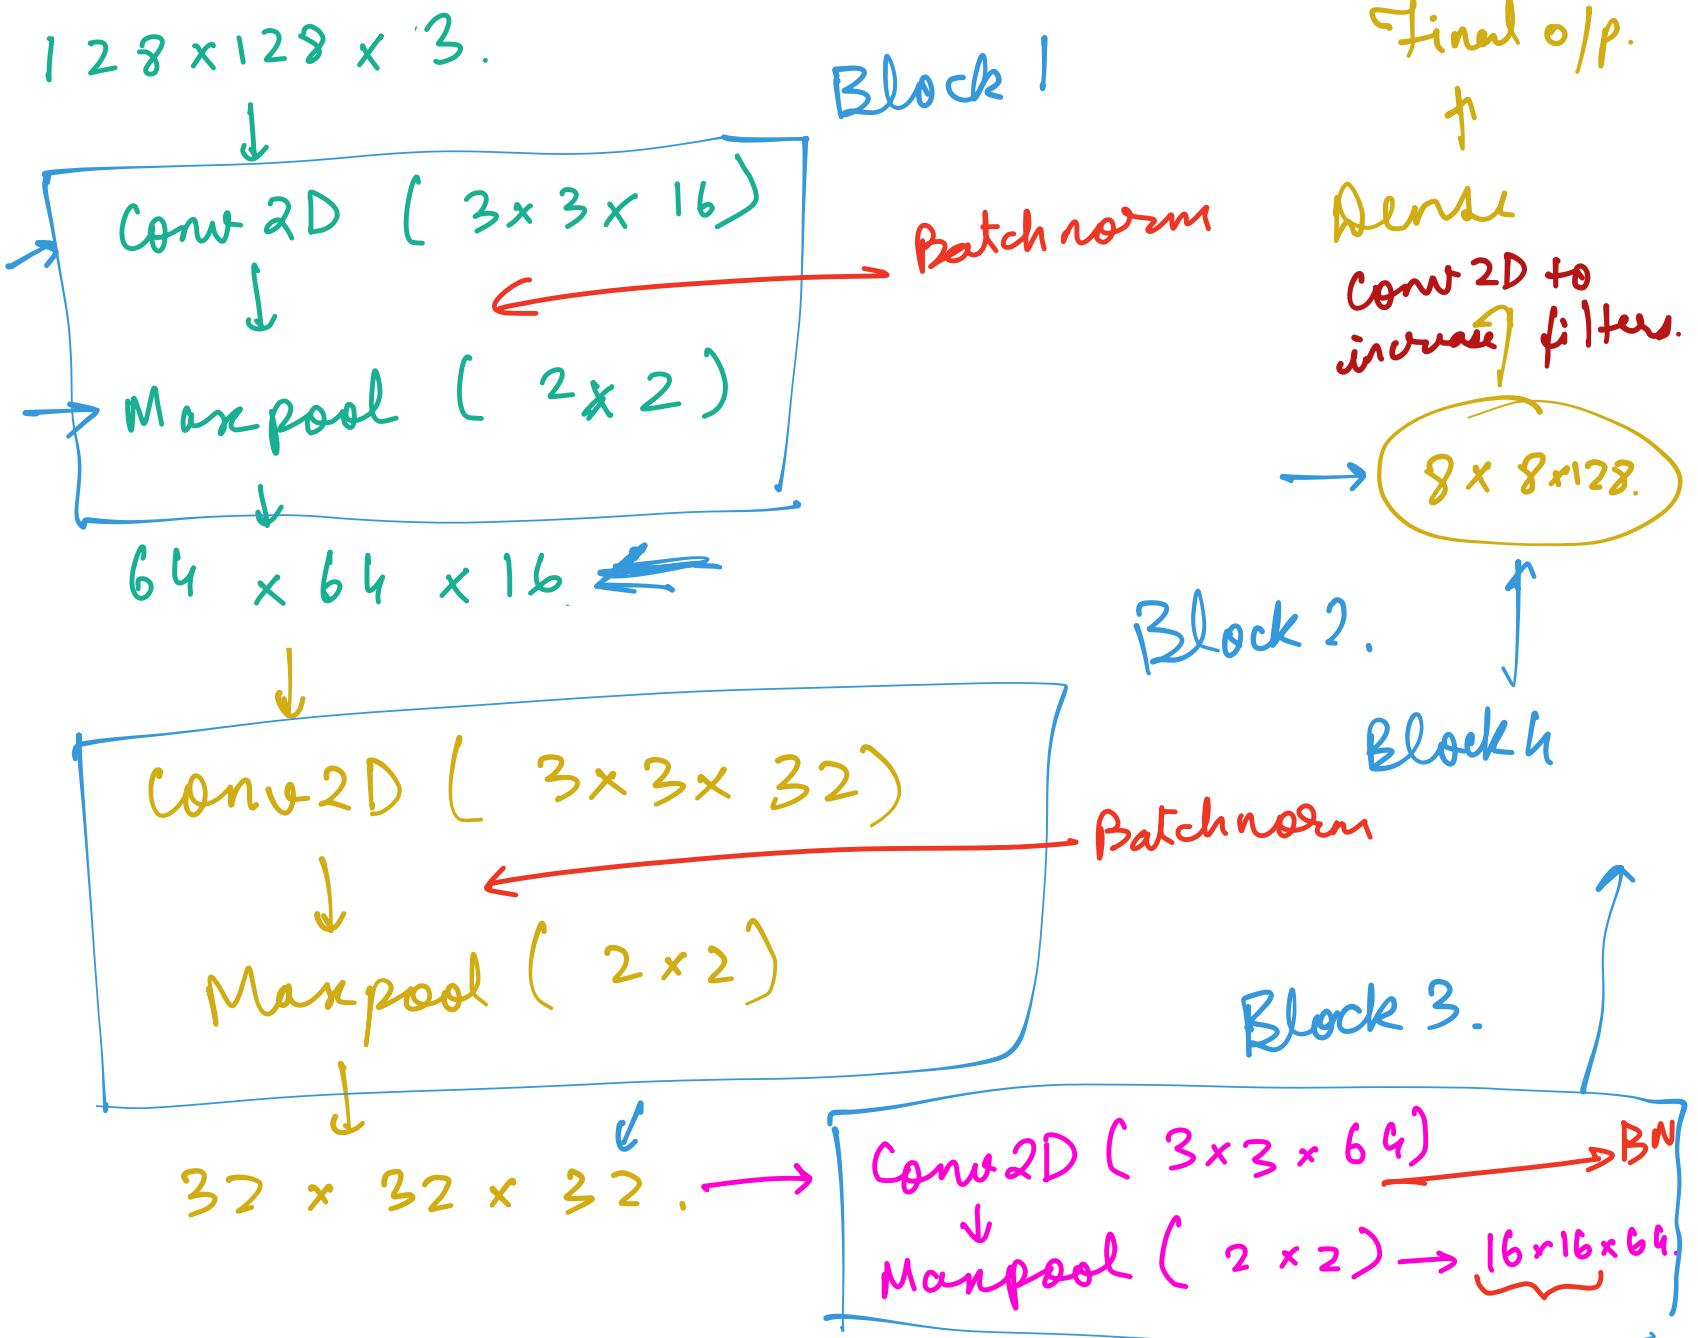
✓ Solution 1 : Maxpool

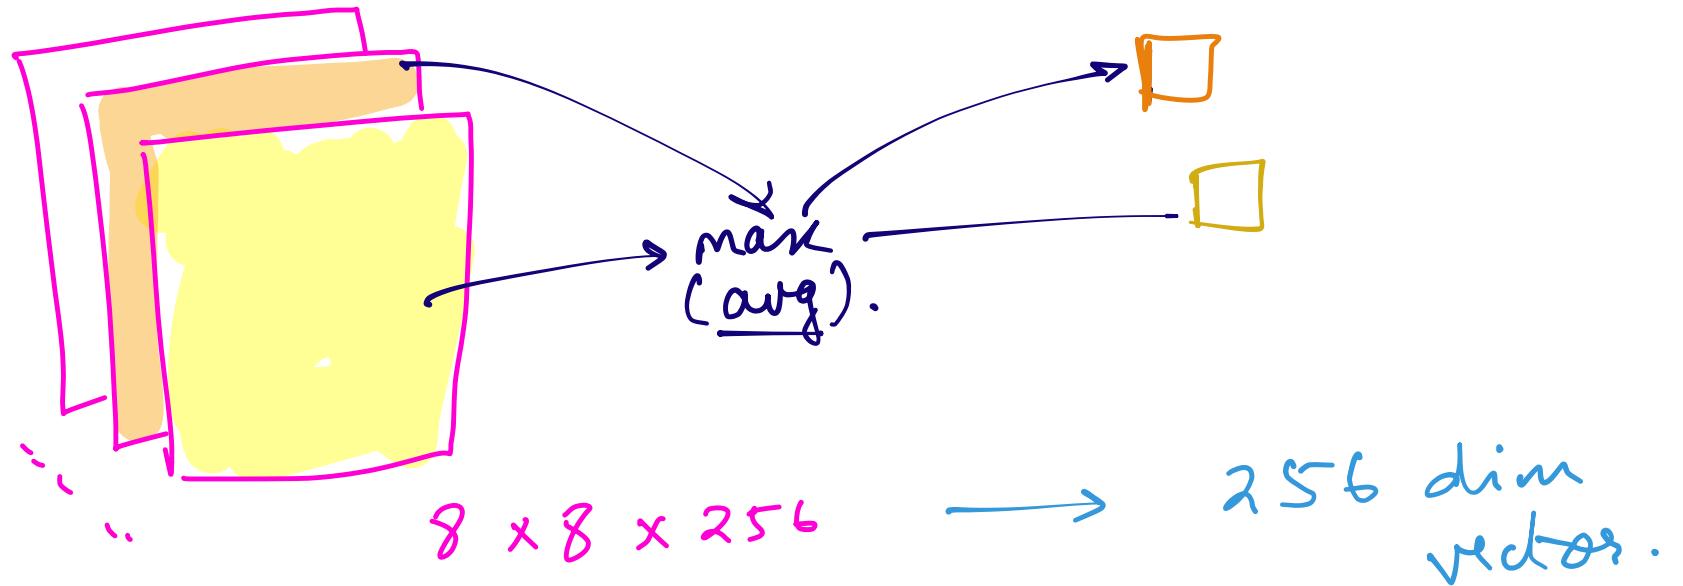
$$n \times n \times k \rightarrow \text{Maxpool2D} \rightarrow \frac{n}{2} \times \frac{n}{2} \times k.$$

✓ Solution 2 : Strided conv.

$$n \times n \times k \rightarrow \text{Conv2D}_{1 \times 1} \rightarrow \frac{n}{2} \times \frac{n}{2} \times k'$$

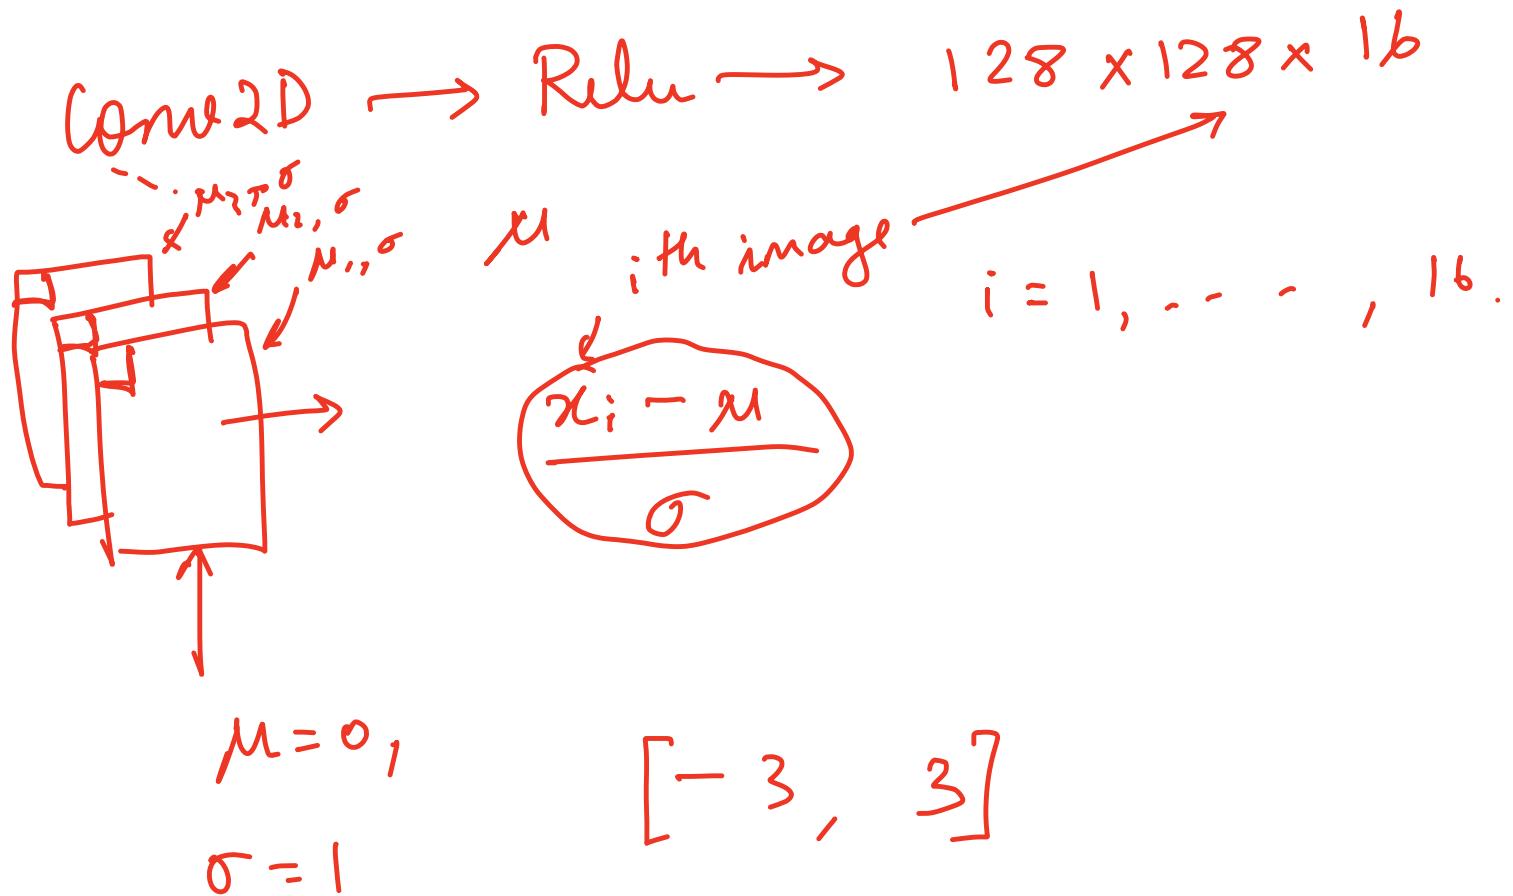
o.p. channels $\rightarrow k'$
stride : 2.

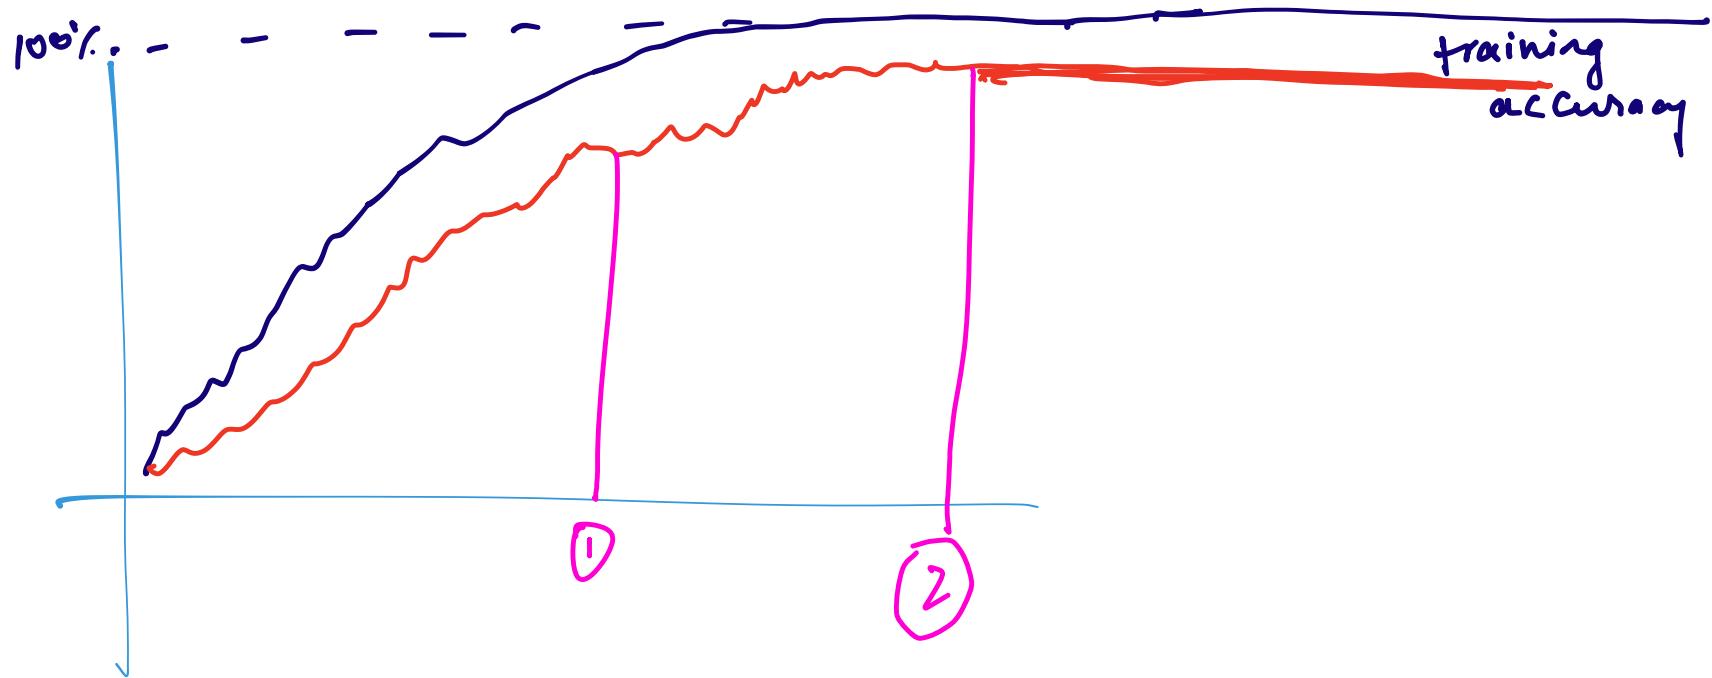




Global Max. pooling
(Avg).

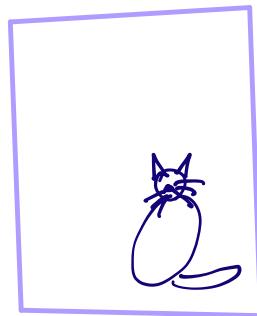
Batch norm



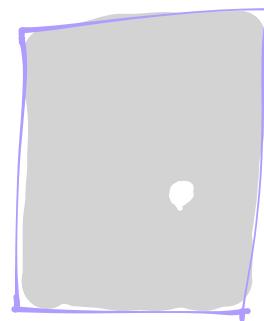


We will next try to stop training when
the Validation accuracy plateaus.

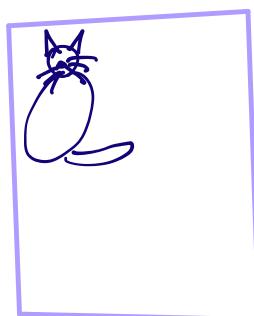
What is translation invariance .



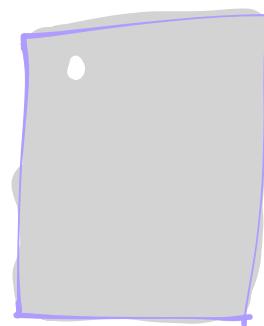
$$* \quad \square =$$



cat face
detector

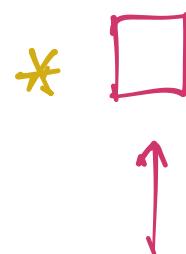
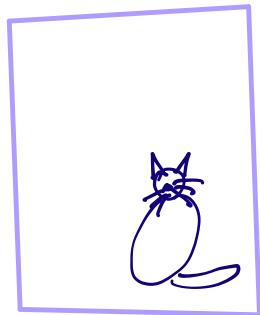


$$* \quad \square =$$



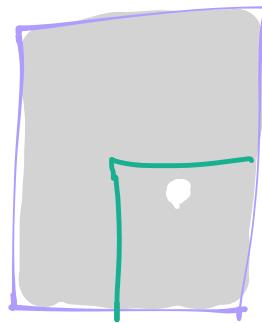
cat face
detector

If we shift the input, and see that the output also shifts, we say that the operation is covariant

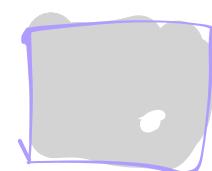


cat face
detector

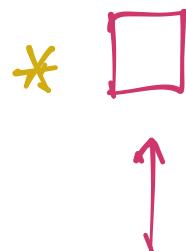
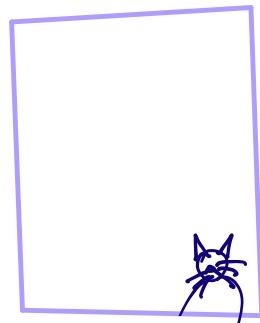
=



Maxpool

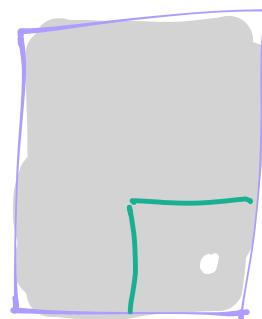


(local) translation
invariance!

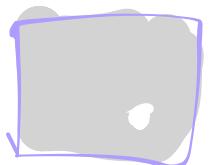


cat face
detector

=



Maxpool



(optional: Deriving gradients for a convolution).

Image $\begin{bmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{bmatrix}_{3 \times 3}$ * $\begin{bmatrix} k_{11} & k_{12} \\ k_{21} & k_{22} \end{bmatrix}_{2 \times 2}$ kernel = $\begin{bmatrix} y_{11} & y_{12} \\ y_{21} & y_{22} \end{bmatrix}_{2 \times 2}$

$$\begin{aligned} y_{11} &= k_{11}x_{11} + k_{12}x_{12} + k_{21}x_{21} + k_{22}x_{22} \\ y_{12} &= k_{11}x_{12} + k_{12}x_{13} + k_{21}x_{22} + k_{22}x_{23} \\ y_{21} &= k_{11}x_{21} + k_{12}x_{22} + k_{21}x_{31} + k_{22}x_{32} \\ y_{22} &= k_{11}x_{22} + k_{12}x_{23} + k_{21}x_{32} + k_{22}x_{33}. \end{aligned}$$

$\frac{\partial d}{\partial k_{11}}$ → what we want.

$$= \left(\frac{\partial d}{\partial Y} \right) \cdot \left(\frac{\partial Y}{\partial k_{11}} \right)$$

$\begin{bmatrix} x_{11} \\ x_{21} \\ x_{31} \end{bmatrix}$ x_{12}
 x_{12} x_{22}